



# The Tech Evolution of Mailinator.com

## Scaling to Millions of Emails/day

---

Paul Tyma, Ph.D.



# What is Mailinator.com?

- A website that accepts email for ANY address @mailinator.com
  - That's ~~billions~~ ~~trillions~~ ~~quadrillions~~ LOTS of possible addresses
- There are no user “accounts”
- Anyone can view any inbox or read any email
- No Signup
- No Login
- Receive-only
- No Attachments
- All email is deleted after a few hours



# What is Mailinator.com?

- As relevant to this talk:
  - A website I started in 2003
  - How it evolved into a SaaS Service with many thousands of daily users including Privacy, Security, Reliability, and Scalability

# What this talk is about

This is not necessarily a talk about the right way to scale a system

This is a talk about how I scaled a system and what I learned along the way

# Who am I

---

- Until Sep 2019, CTO @lendingtree.com
- Founded 5 startups
  - Preemptive Solutions, Pulse.io, Home-account, Manybrain (Mailinator), Refresh
- A few years at Google, Linked
- Ph.D. in Computer Engineering, 2004 - ~compilers

# Disclaimer

---

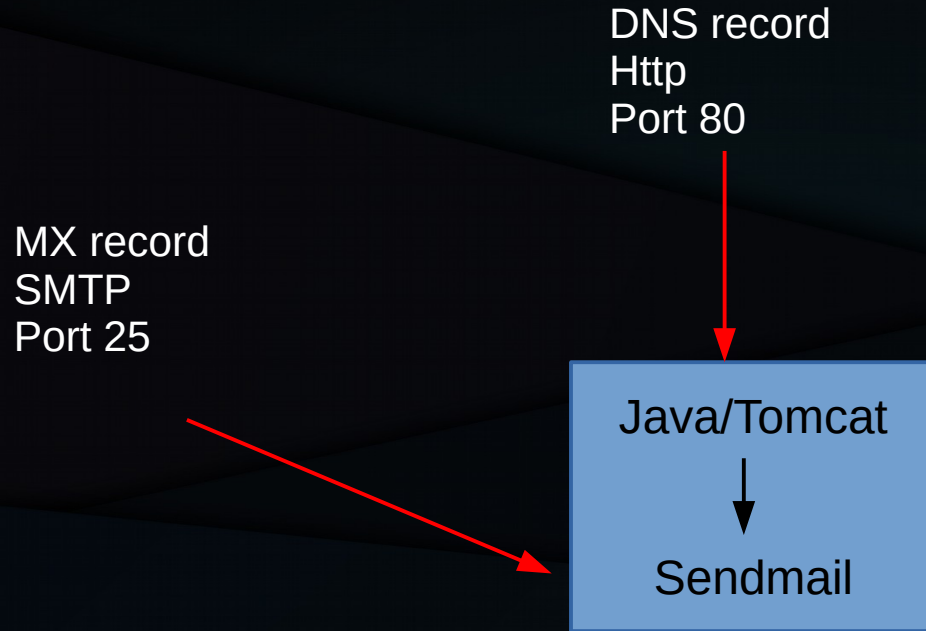
- We'll be showing live (text-only) data
- We don't control the data
- There might be bad words in the data
- There might be a **lot** of bad words in the data
- Only click emails you're sure you want to see

# Initial implementation over a weekend

- Configure Sendmail:
  - Vim virtusertable  
@mailinator.com catchallinbox
- Java (+imap library) / Tomcat
- Code up terrible UI
- Rent Serverbeach:
  - AMD 2Ghx Athlon, 1G ram, 80G HD, 10Mbps Net
- Store only the 20,000 most recent emails (FIFO)



# Hardware



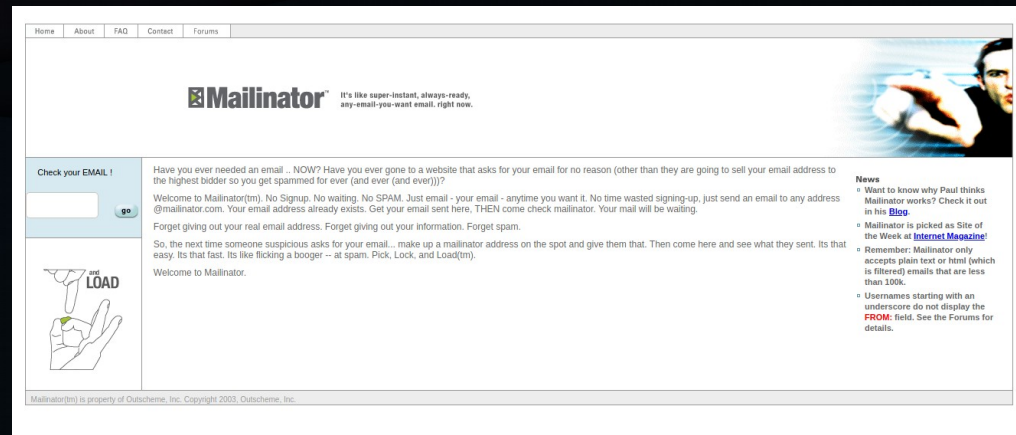


# Launch !

- 6000 web visitors first day thanks to Joel

July 23, 2003 by Joel Spolsky

Need a disposable email address so you can sign up for some web site, quick? You've already got one, at Paul Tyma's clever Mailinator. "Just send an email to any address @mailinator.com. Your email address already exists. Get your email sent here, THEN come check mailinator. Your mail will be waiting." More proof that great UI design is done by taking away, not adding things.



# Things I learned

- The world tends to think an email address implies identity
  - The world is wrong
- Some Websites that allow signup with an email address figure that creating an email address takes time and that will act as a natural rate-limiter
  - Those websites are wrong
- We write lots of anti-abuse, rate-limiting code

# System Limitations

- Both good and bad Software Engineers build systems that eventually crash under load
  - But the good ones can tell you where it'll happen first
- Disk seek time (in early 2000's): ~10ms
- Each incoming email requires a disk write and a disk delete (of the pushed out email)
  - Hand waving over many other disk seeks (directory structures, etc)
- The hard drive alone limited us to ~50 emails/second
- Disk cache was useless given data ingestion pattern

# System Limitations (circa 2005)

- System crashed on the regular at about 800,000 emails per day (~9.3/second)
- User's reading email caused further disk reads but a trivial amount in the big picture
- As is often the case, disk was the problem
  - Note: SSD seeks times now around ~65microseconds
  - Other Note: Memory access measured in nanoseconds

# Rewrite to store data in RAM (circa 2005)

- Google keeps the entire web index in (lots of) RAM
- Why not store all email in RAM?
  - 20,000 emails \* ~4k email size = 80M
  - Plus Index inbox lists and email id's
  - Java isn't known for memory frugality
  - Estimate 130M per 20,000 emails



# Rewrite to store data in RAM (circa 2005)

- Build a custom Multithreaded SMTP server
  - SMTP is a very chatty protocol
  - thread-per-connection
  - Up to 300 simultaneous threads
    - 512k stack per thread
- Upgrade server to 2GB

```
S: 220 mailserver.example.com ESMTP Postfix
C: HELO sender.test.com
S: 250 Hello mailsender.test.com, I am glad to meet you
C: MAIL FROM:alice@test.com
S: 250 Ok
C: RCPT TO:bob@example.com
S: 250 Ok
C: RCPT TO:john@example.com
S: 250 Ok
C: DATA
S: 354 End data with <CR><LF>.<CR><LF>
C: From: "Alice Test" alice@test.com
C: To: Bob Example bob@example.com
C: Cc: john@example.com
C: Date: Sun, 16 Feb 2014 18:04:25 +0530
C: Subject: Test message
C:
C: Hello Bob.
C: Sending a test message to test the SMTP protocol operation.
C: Yours Sincerely ,
C: Alice
C:
S: 250 Ok: queued as 12345
C: QUITS: 221 Bye
(The server closes the connection)
```

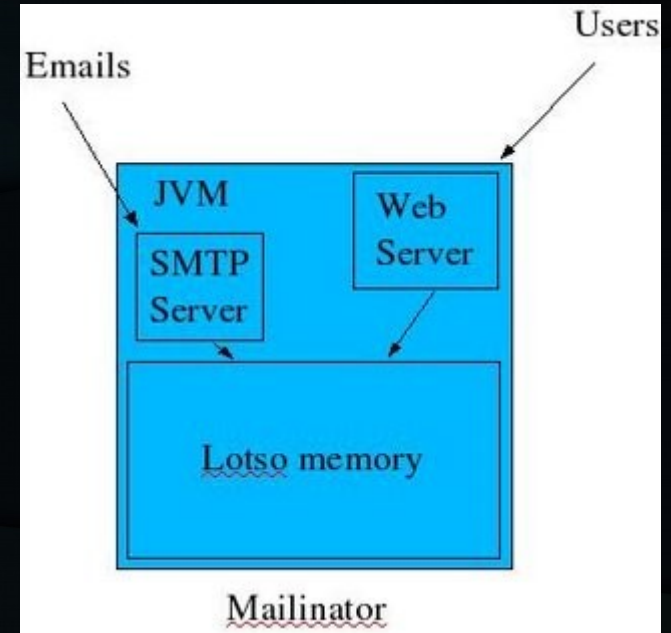
Initial Handshake

Message Header

Message Body

# Rewrite to store data in RAM (circa 2005)

- Web + SMTP in one
  - Ports 25, 587, 80, 443
- Bouncing the server not only brought the site down, it now lost ALL email
  - Hey. It's a free service.





# How do we store these emails?

- Compactly. RAM is limited
- Must keep ordering of receipt to expire old emails
- Must be able to lookup emails by inbox
- Must be able to retrieve emails by ID

# Storage: If this were a database

# Index



# Index

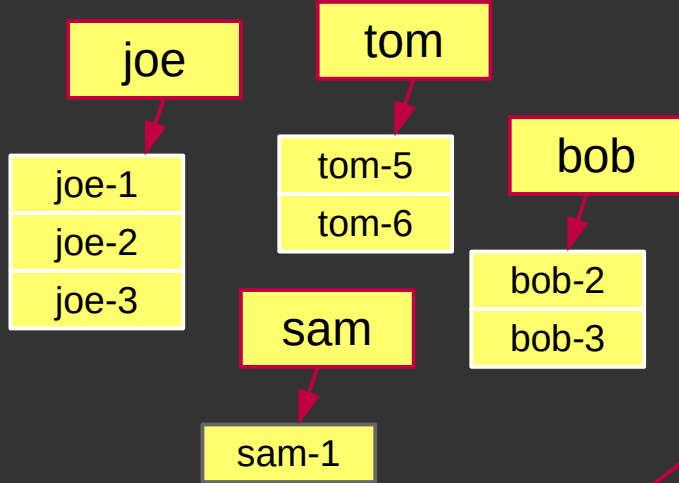


ID	Received	To	From	Subject	Email Body
joe-3	1/12 3:24pm	joe	elon@tesla.com	Gasoline sucks	<pre>-----_NextPart_001_000_010E8AB_A6AFAF0 Content-Type: text/plain; charset=UTF-8 Content-Transfer-Encoding: base64  This is HTML  -----_NextPart_001_000_010E8AB_A6AFAF0 Content-Type: text/html; charset=UTF-8 Content-Transfer-Encoding: quoted-printable  &lt;html&gt;&lt;body&gt;&lt;div id=</pre>
bob-5	1/12 3:25pm	bob	billg@microsoft.com	Try windows	<pre>-----_NextPart_001_000_010E8AB_A6AFAF0 Content-Type: text/plain; charset=UTF-8 Content-Transfer-Encoding: base64  This is HTML  -----_NextPart_001_000_010E8AB_A6AFAF0 Content-Type: text/html; charset=UTF-8 Content-Transfer-Encoding: quoted-printable  &lt;html&gt;&lt;body&gt;&lt;div id=</pre>
frank-9	1/12 3:26pm	frank	jeff@amazon.com	Buy now	<pre>-----_NextPart_001_000_010E8AB_A6AFAF0 Content-Type: text/plain; charset=UTF-8 Content-Transfer-Encoding: base64  This is HTML  -----_NextPart_001_000_010E8AB_A6AFAF0 Content-Type: text/html; charset=UTF-8 Content-Transfer-Encoding: quoted-printable  &lt;html&gt;&lt;body&gt;&lt;div id=</pre>

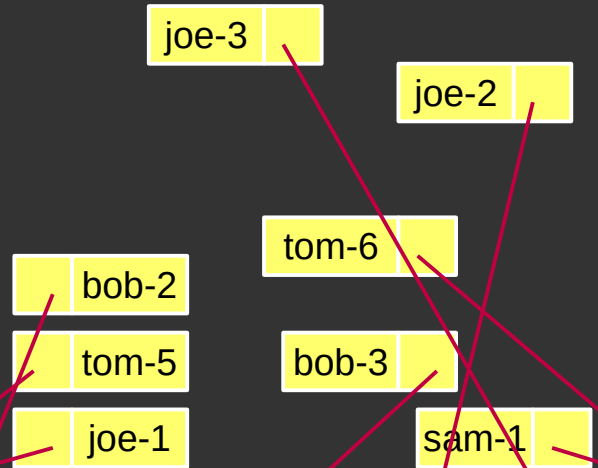


# How to Organize Memory (V1)

## Inbox Dictionary/HashMap

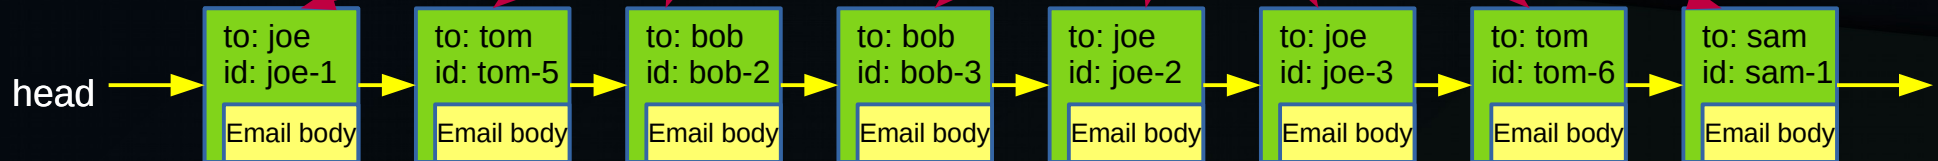


## Email Dictionary/Hashmap



Fetch Inbox: 100's per minute  
Fetch Email: 10's per minute  
Add Email: 1000's per second  
(User) Delete Email: 1's per minute

Need mutex for all 3 structs



# Organize Memory

## Considerations:

- Inbox Dictionary: hundreds of thousands of entries
  - EmailId Dictionary: millions of entries
  - Email LinkedList: millions of entries
  - Unlike a lot of Dictionary usage, we have a 10/90 read/write ratio
- 
- All 3 structs must be synchronized for mutation
    - i.e. every incoming email
    - Dictionary resizing holds locks a long time
  - Replace all structs with: SynchronizedLinkedTreeMap

# How about compressing emails?

- CPU intensive
- Became feasible with higher-core machines
- Doesn't work well
  - Compression algorithms typically work by building a dictionary. Emails are often too short to build a notable one
  - Base64 emails (i.e. big emails) build none at all
  - Tried a email-corpus Huffman encoder
    - Worked on Headers
    - Body data not so much

# How about reusing emails?

Received: from mail-ot1-f43.google.com([209.85.210.43])  
by mail.mailinator.com with SMTP (Postfix)  
for joe@mailinator.com;  
Mon, 06 Jan 2020 20:54:52 +0000 (UTC)  
From: Paul Tyma <paul@manybrain.com>  
Date: Mon, 6 Jan 2020 15:54:41 -0500  
Subject: Hi Joe  
To: joe@mailinator.com  
Content-Type: multipart/alternative; boundary="000000000000ec3632059b7ede21"

headers

--000000000000ec3632059b7ede21  
Content-Type: text/plain; charset="UTF-8"

Hey \*Joe\*. How are you today? Buy our debt consolidation services !!!!

--000000000000ec3632059b7ede21  
Content-Type: text/html; charset="UTF-8"

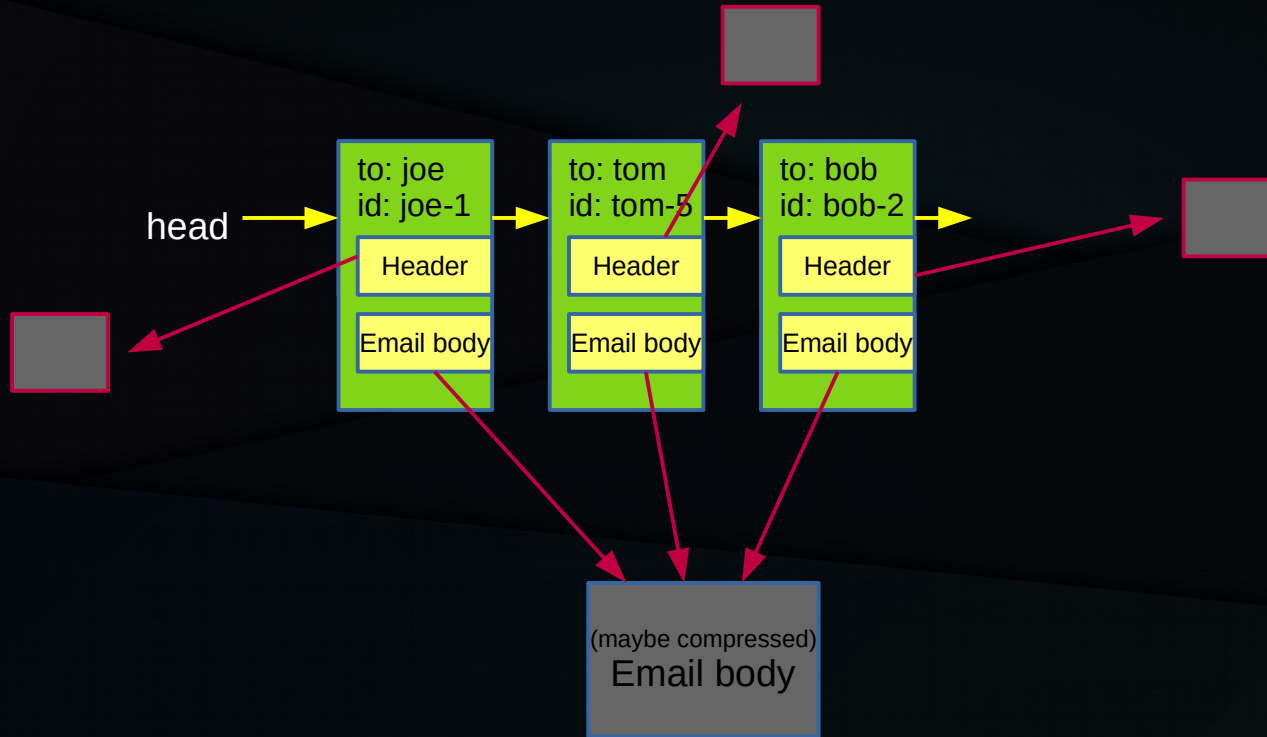
<div dir="ltr"><div class="gmail\_default" style="font-family:tahoma,sans-serif;font-size:large">Hey <b>Joe</b>.</div></div>

--000000000000ec3632059b7ede21--

body



# Compress and Reuse Message bodies





# Everything is Great!

- New System Launch !!!
- System went from crashing at 800,000 to running fine at a consistent 1MM emails/day
  - Rock solid and super fast!
- Disk completely idle
- CPU's not busy at all

# Or is it?

---

- Kept getting around ~1MM emails/day
- Went on like this for months - it became “too” consistent
- Site seemed pretty slow
- Customers kept complaining email didn’t arrive
  - What did they know

# There's always a new bottleneck

- Then - Serverbeach had a sale
  - Upgraded the 10Mbps NIC to 100Mbps NIC
- Traffic immediately shot up to 3MM emails/day
  - 10Mbps = 1.25MB/s
  - 100Mbps = 12.5MB/s
- Every system has a bottleneck
  - Choose: CPU, RAM, Disk, Network

# Things I learned:

- Every user (or email) that hits your system is part of a concerted Denial-of-service attack
  - It's much harder to DoS your system if their request only hits RAM and CPU
  - It's really easy to DoS your system if their request hits disk

# Things I learned:

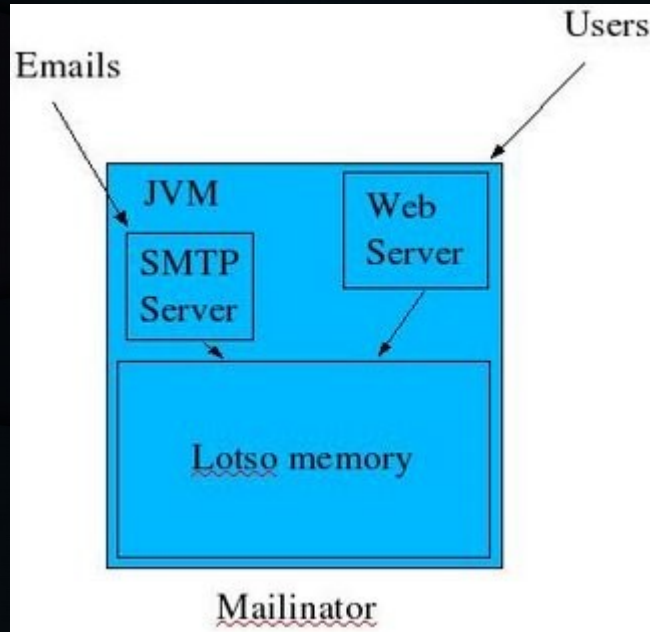
- Forgot to check domain on incoming emails
  - We don't just accept email for any inbox @mailinator.com – we accept incoming email for any inbox at **any domain** that hits our server
    - **That's a Feature !**
    - Today, ~1900 domains point their MX record at us
    - <https://viewdns.info/reversemx/?mx=mail.mailinator.com>
  - That looks like an open relay
    - We get a lot email from security people telling us our email server is miconfigured
    - Email FROM @mailinator.com always gets classified as Spam
    - But we don't send email?
    - And we don't want to !
    - **Another Feature !**

# Everything is Great!

- But still running on a single server
  - Few million emails per day
  - Few thousand users per day
- Then I went on vacation for a week
  - Hardware failures always happen when you're on vacation
  - Need to make Infrastructure redundant



# Hardware (where we are)



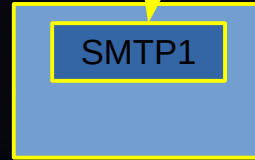
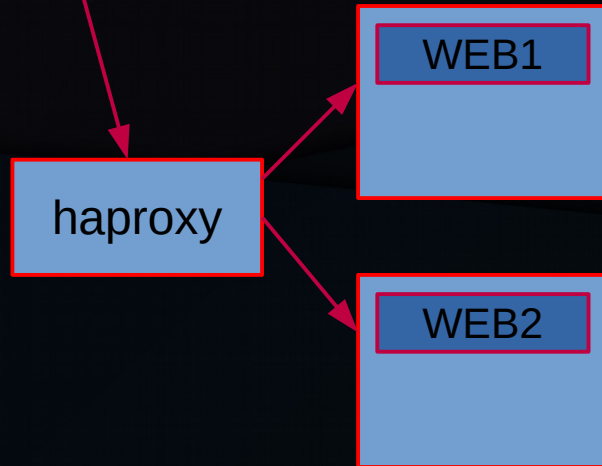


# Scaling Thought Process

WEB  
HTTPS:443

Public Email  
MX1:25

Public Email  
MX2:25



Remember:

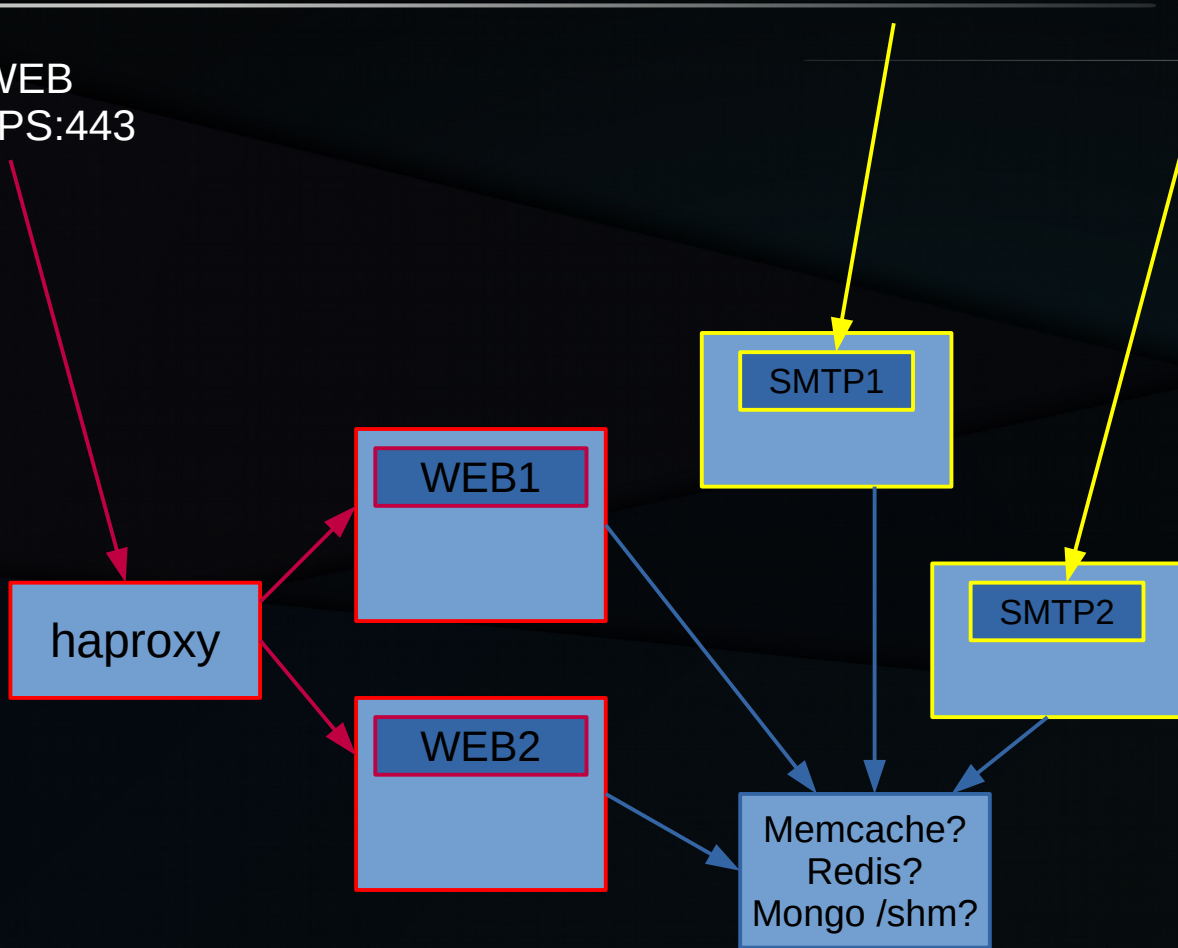
- Each Rectangle is a physical (i.e. virtual) server
- We get WAY more SMTP traffic than Web traffic (but web is prioritized)

# Central Store?

WEB  
HTTPS:443

Public Email  
MX1:25

Public Email  
MX2:25



So:

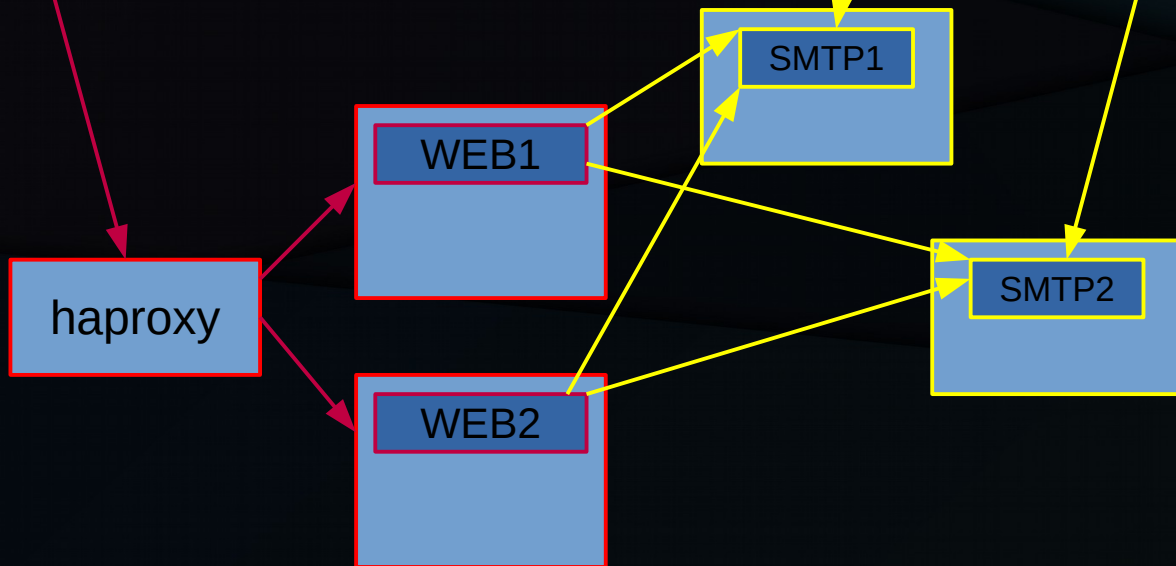
- Central store introduces a bottleneck
- SMTP servers have doubled their bandwidth
- Like most systems in the world, the datastore is surely to be our next bottleneck
- Horizontal scaling of SMTP servers does not help us

# Scalability

WEB  
HTTPS:443

Public Email  
MX1:25

Public Email  
MX2:25



So:

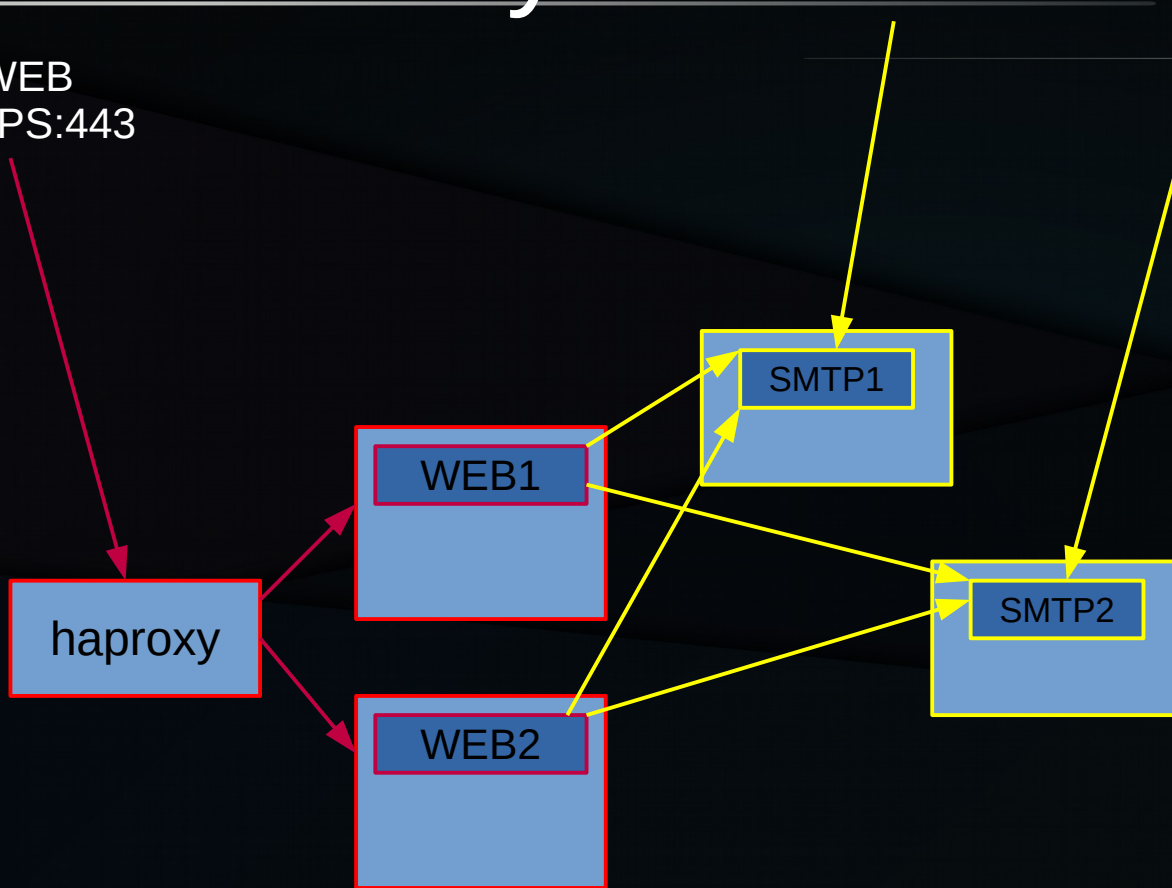
- What happens in the SMTP server, Stays in the SMTP server
- MX records scale horizontally
- SMTP Servers don't know about each other – we can continue to add more of these
- Web servers unloaded and Cloudflare protects them from shenanigans
- Can scale them too up to HAPROXY limitations (which is really high)

# Redundancy

WEB  
HTTPS:443

Public Email  
MX1:25

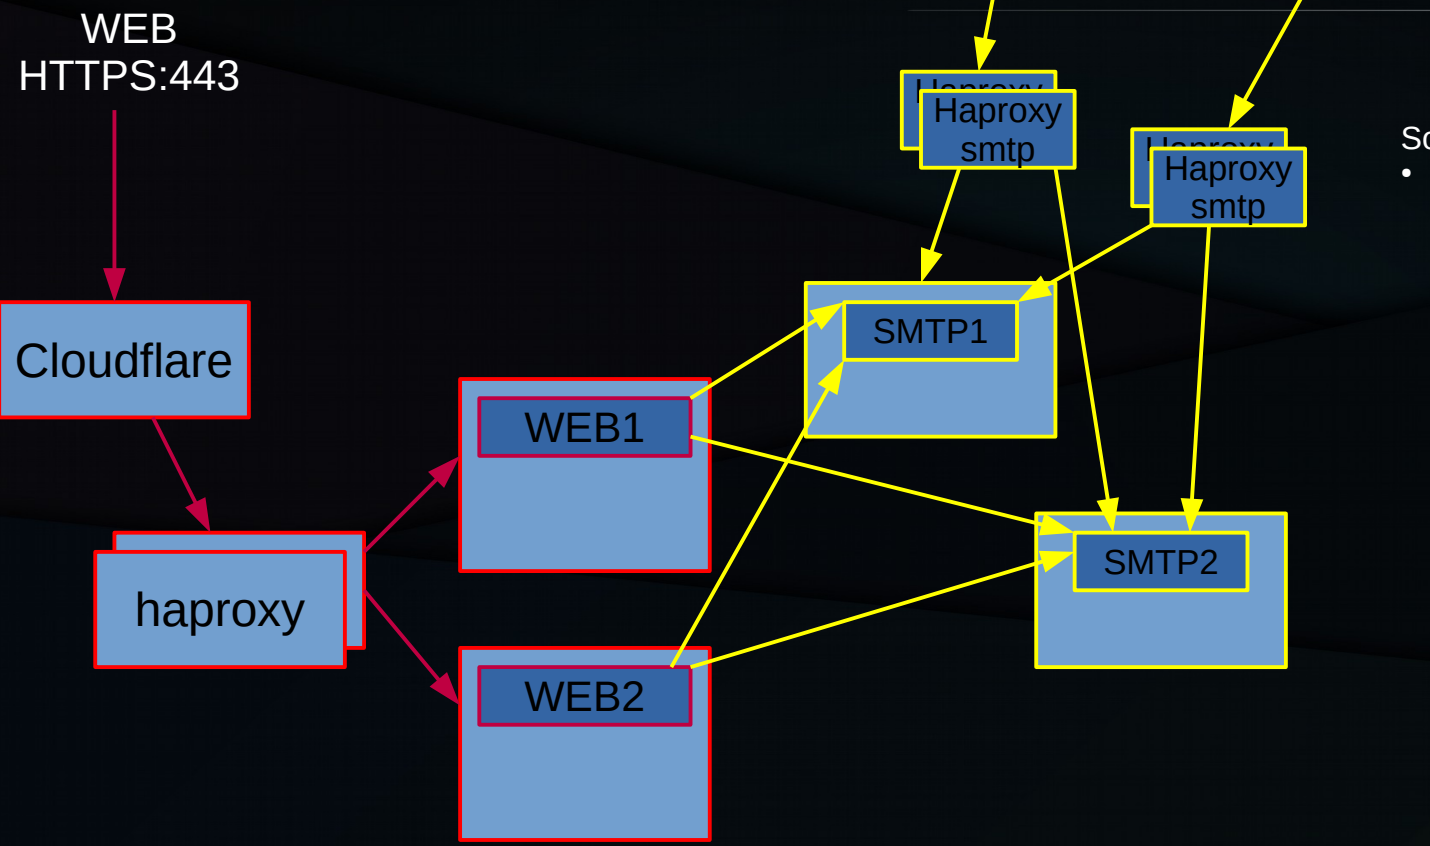
Public Email  
MX2:25



So:

- If a webserver OR Smtptserver goes down. We're fine.
- Web talks to SMTP through proprietary socket protocol (not REST). Note that pure TCP connections have a 3 way SYN, SYN-ACK, ACK just to connect.
- Haproxy never changes and is super stable.
- MX records naturally distribute the load between SMTP servers. But sadly, not very well.
- If a SMTP server goes down, we're still not receiving some email
- We might get DDos'd (hint: we did). Let's get behind a Ddos protector.

# Reliability

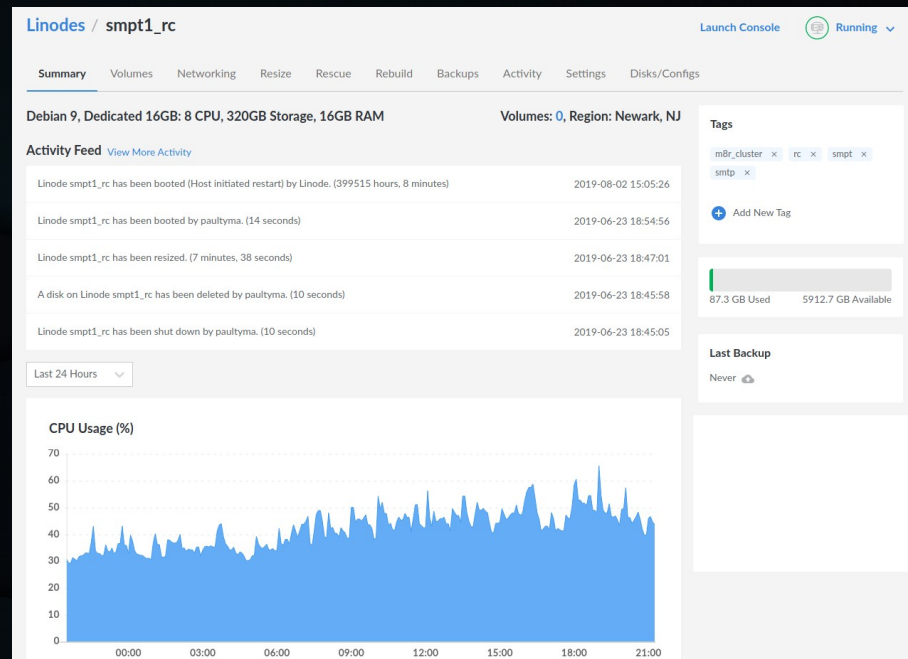


So:

- Great! Super stable. Knock down anything and we're fine. Reboot anything and we're fine. (except Cloudflare)

# Multi-server problems

- Service Discovery
  - Linode API with Tags
- Jgroups
  - Leader Election
    - Batch Jobs!
  - Fabric
    - Cross-server shared heap





# Everything is Great !

- System running like clockwork
- Only thing that brings system down is me tinkering with the code
- Support emails start to increase
  - But not for system issues
  - Mostly business users asking for features and enhancements



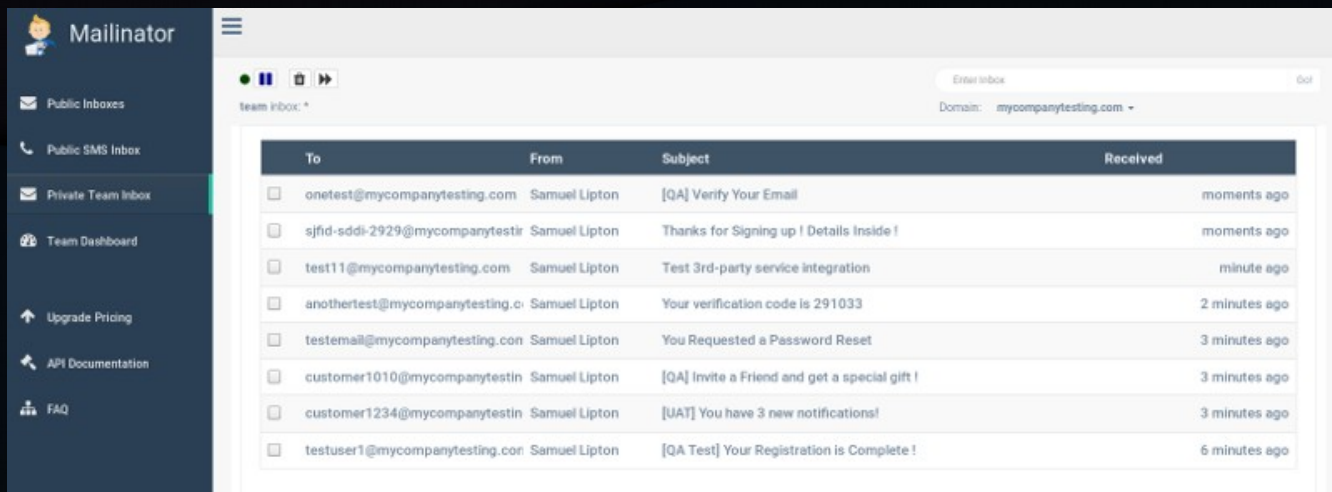
# Time to go Pro

---

- You know - We could charge for this
- QA Teams are using this like crazy
  - Having an infinite number of email addresses is darn handy for testing

# Subscriber Features

- Your own “Private Domain” - i.e. @mycompanytesting.com
  - Email is private to your team
  - Private emails aren’t stored in RAM – they stick around as long as you need them
  - Web Interface now shows all inboxes for an **entire domain**
  - Search inboxes with wildcards test\*@mycompanydomain.com



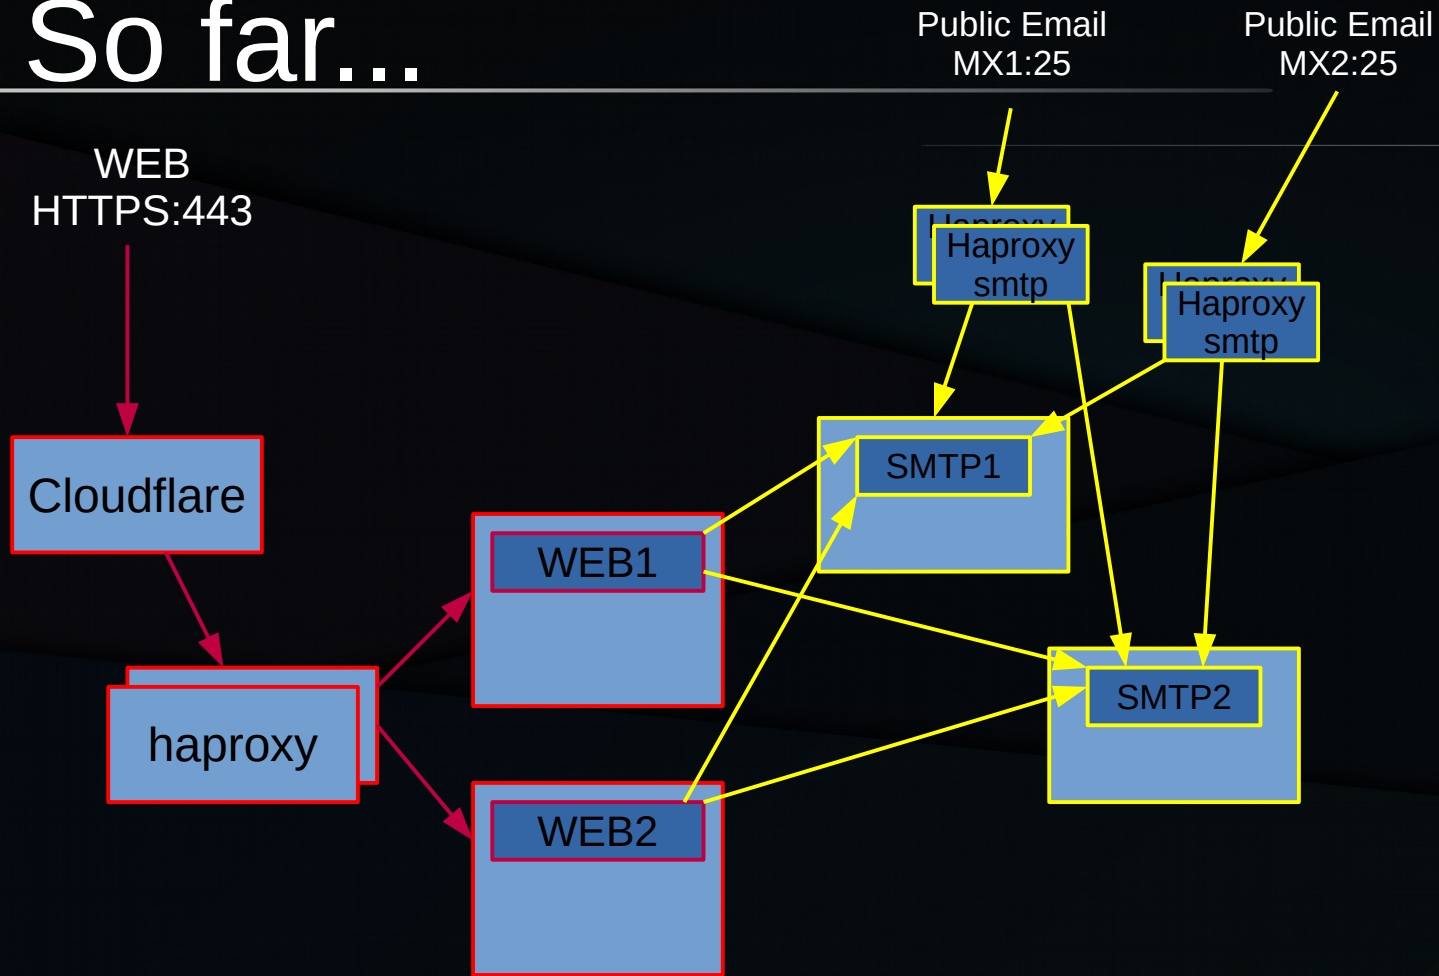
The screenshot displays the Mailinator web interface. On the left is a dark blue sidebar with the Mailinator logo and navigation links: Public Inboxes, Public SMS Inbox, Private Team Inbox (highlighted), Team Dashboard, Upgrade Pricing, API Documentation, and FAQ. The main content area shows a list of emails for the domain 'mycompanytesting.com'. The list has columns for 'To', 'From', 'Subject', and 'Received'. Each email entry includes a checkbox, the recipient email address, the sender 'Samuel Lipton', the subject line, and the time received.

	To	From	Subject	Received
<input type="checkbox"/>	onetest@mycompanytesting.com	Samuel Lipton	[QA] Verify Your Email	moments ago
<input type="checkbox"/>	sjfid-sddi-2929@mycompanytestir	Samuel Lipton	Thanks for Signing up ! Details Inside !	moments ago
<input type="checkbox"/>	test11@mycompanytesting.com	Samuel Lipton	Test 3rd-party service integration	minute ago
<input type="checkbox"/>	anotherest@mycompanytesting.co	Samuel Lipton	Your verification code is 291033	2 minutes ago
<input type="checkbox"/>	testemail@mycompanytesting.com	Samuel Lipton	You Requested a Password Reset	3 minutes ago
<input type="checkbox"/>	customer1010@mycompanytestin	Samuel Lipton	[QA] Invite a Friend and get a special gift !	3 minutes ago
<input type="checkbox"/>	customer1234@mycompanytestin	Samuel Lipton	[UAT] You have 3 new notifications!	3 minutes ago
<input type="checkbox"/>	testuser1@mycompanytesting.com	Samuel Lipton	[QA Test] Your Registration is Complete !	6 minutes ago

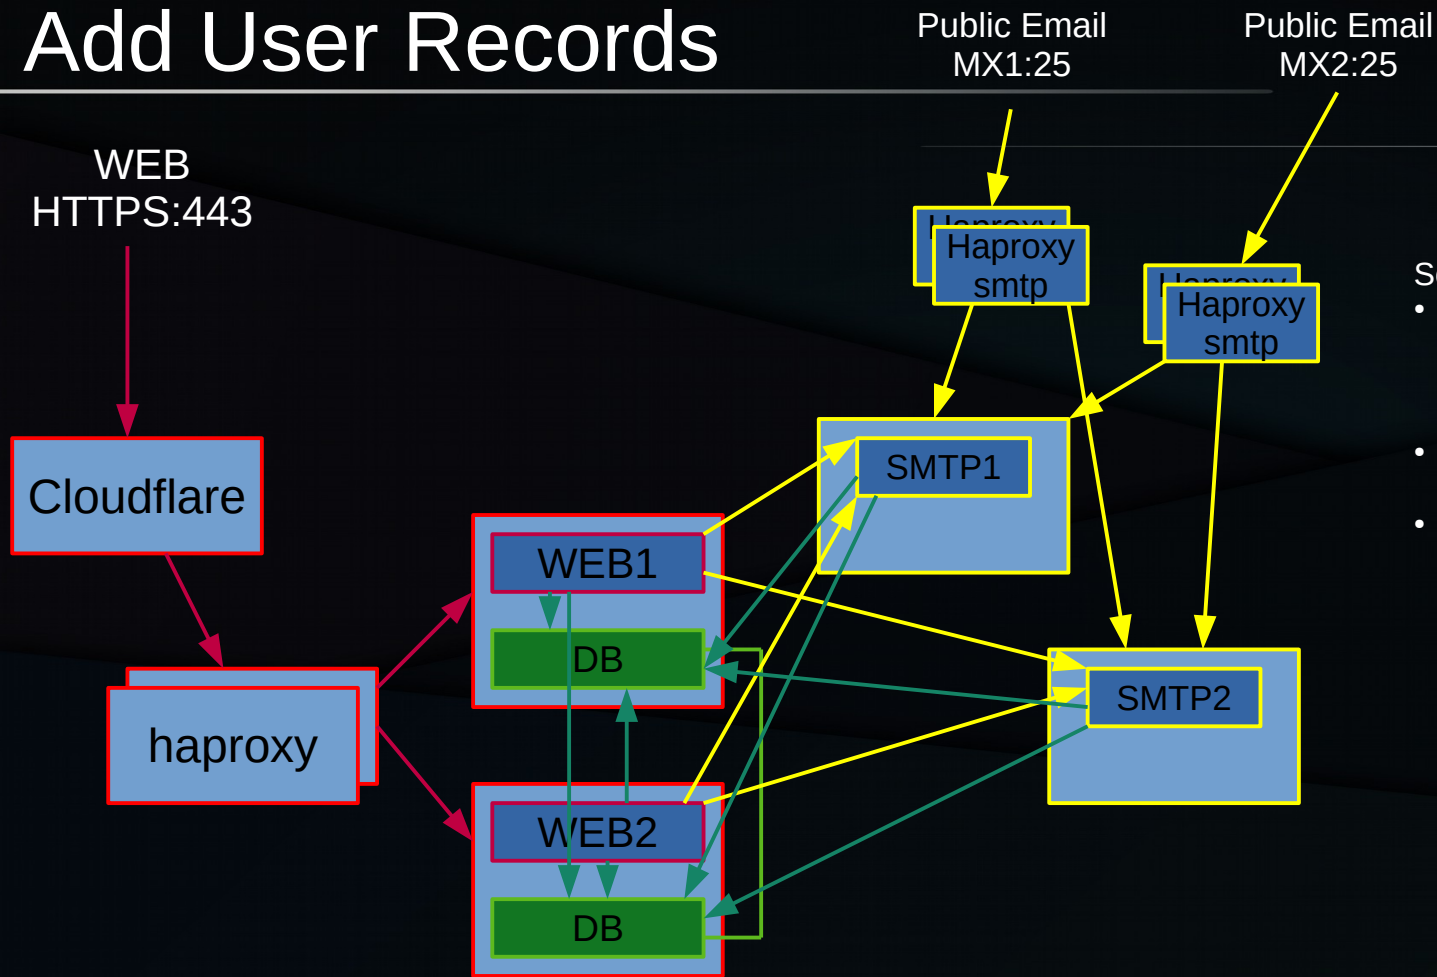
# Subscriber Features

- API Access
- Private SMS testing numbers
  - Incoming SMS's arrive in “inboxes” of the phone number
- Rule System
  - Auto-route emails. Send them to Slack. Forward them. Auto-clicking of links.
- We need SignUp, and Login, and “forgot your password”, and Uptime guarantees, Security, and Support
  - Sheesh... free services are way easier

# So far...



# Add User Records



So:

- Add a DB cluster (redundant) for User records – signup, private domains, etc.
- Low traffic. Just logins and such.
- Connect it to, well, everything

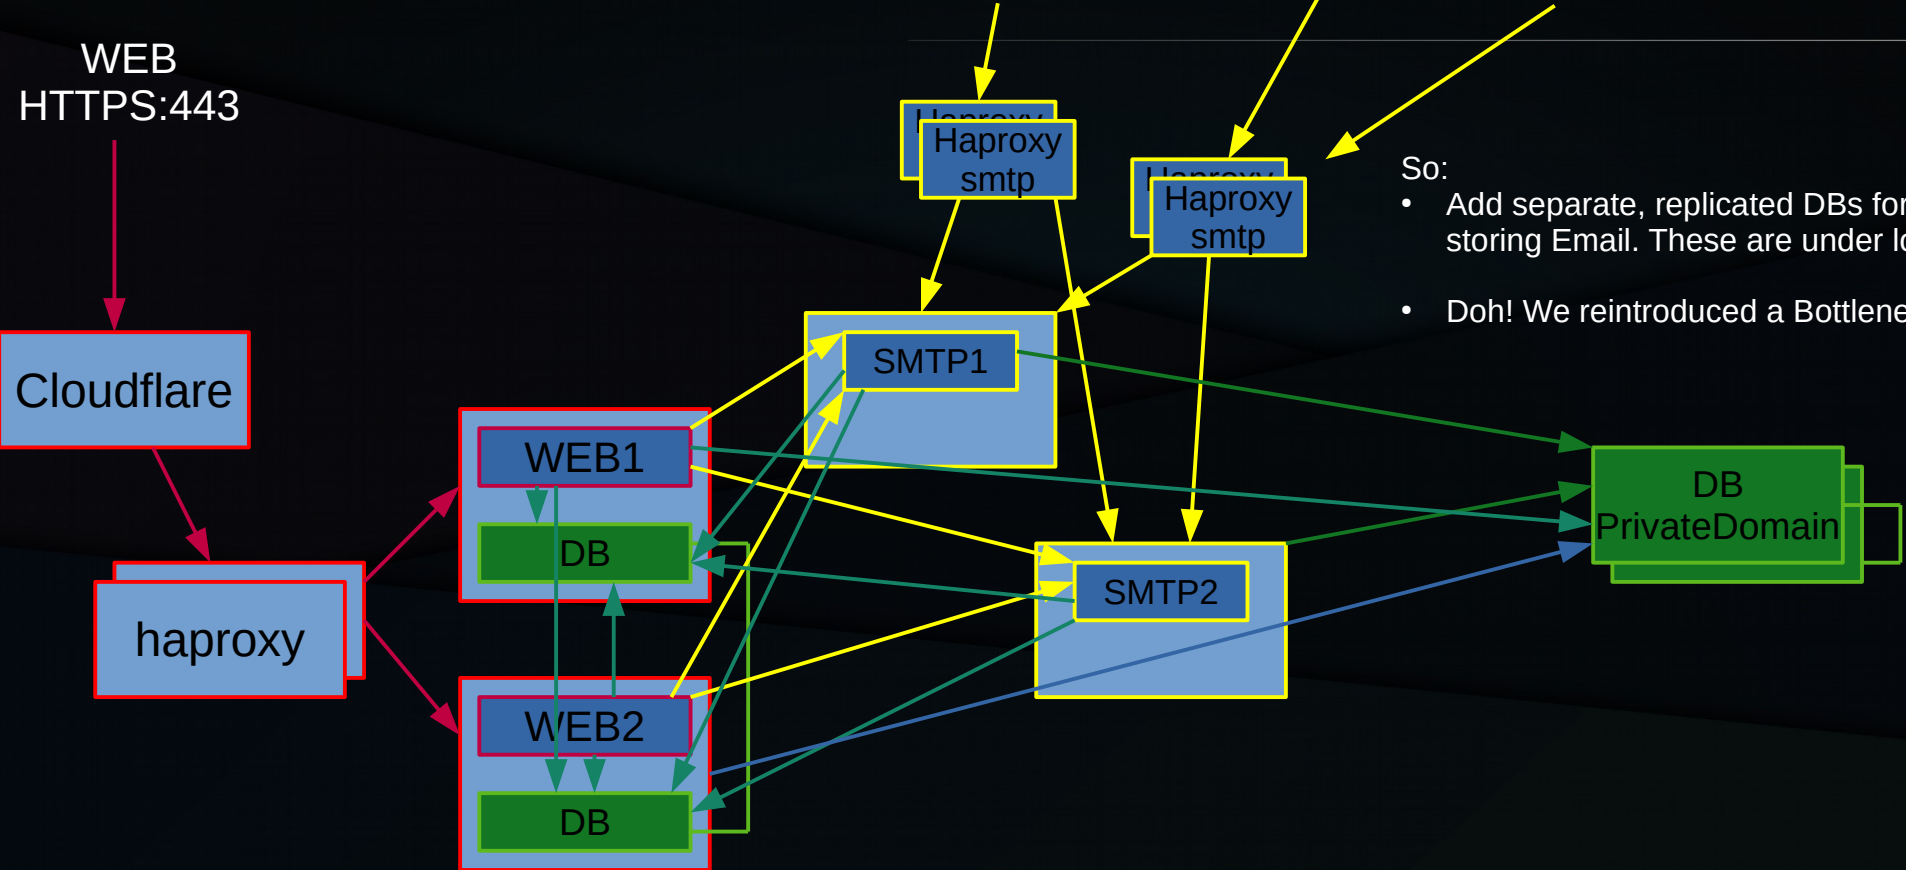
# Private Domains

Public Email  
MX1:25

Public Email  
MX2:25

mail.privatedomain.com  
MX:25

WEB  
HTTPS:443

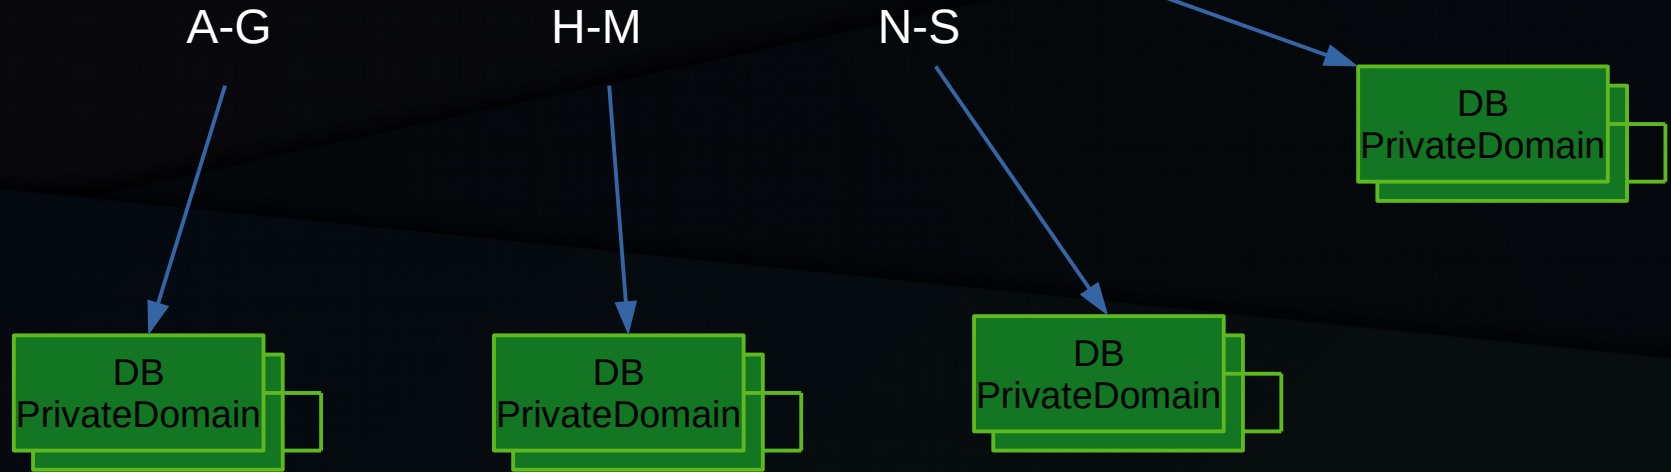


- So:
- Add separate, replicated DBs for storing Email. These are under load
  - Doh! We reintroduced a Bottleneck?



# We Can Shard

Customer Last Name

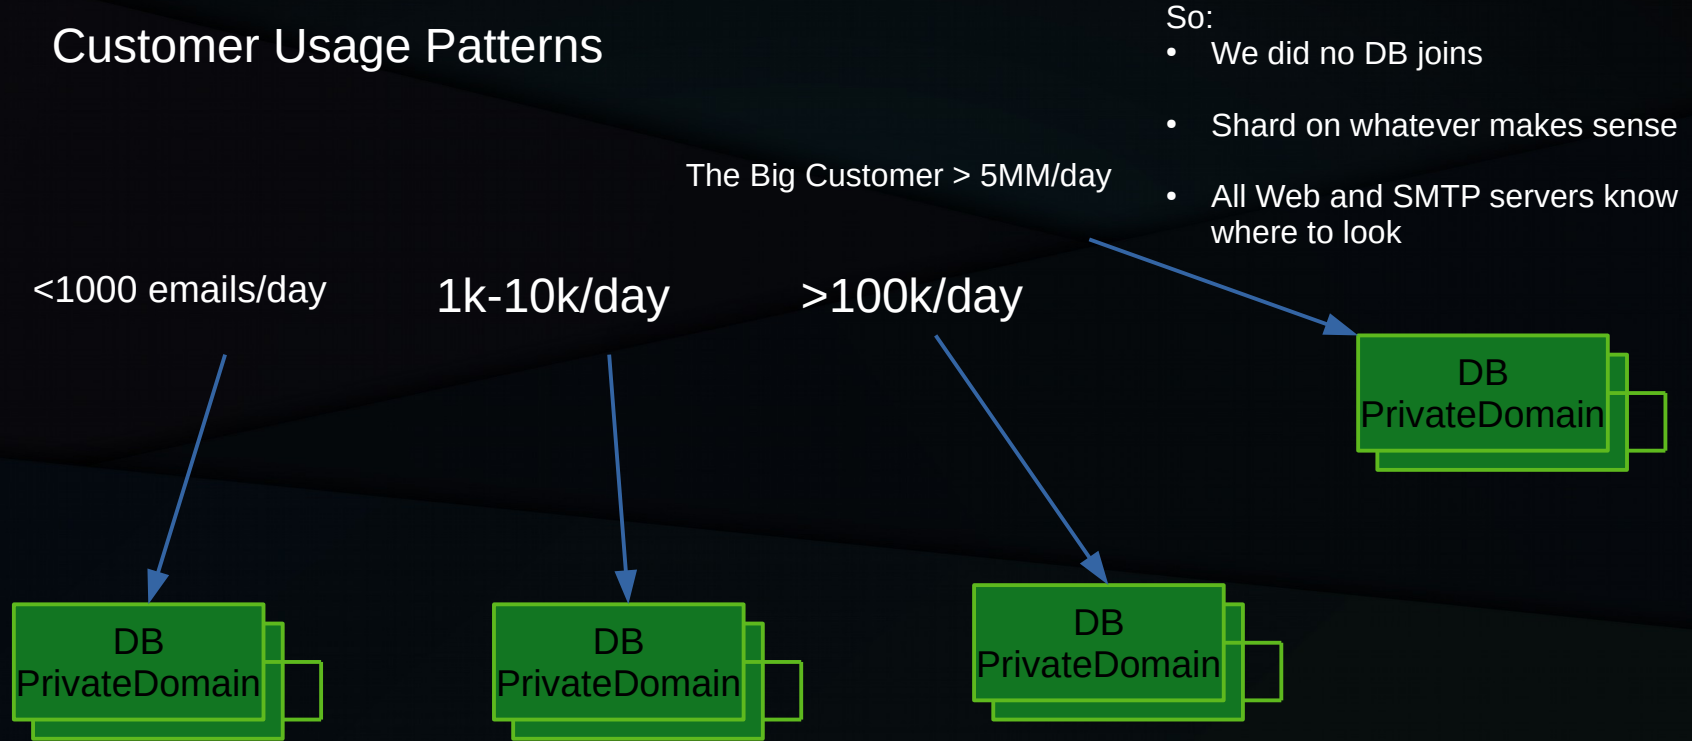


So:

- We did no DB joins
- Shard on whatever makes sense
- All Web and SMTP servers know where to look

# We Can Shard

## Customer Usage Patterns



So:

- We did no DB joins
- Shard on whatever makes sense
- All Web and SMTP servers know where to look

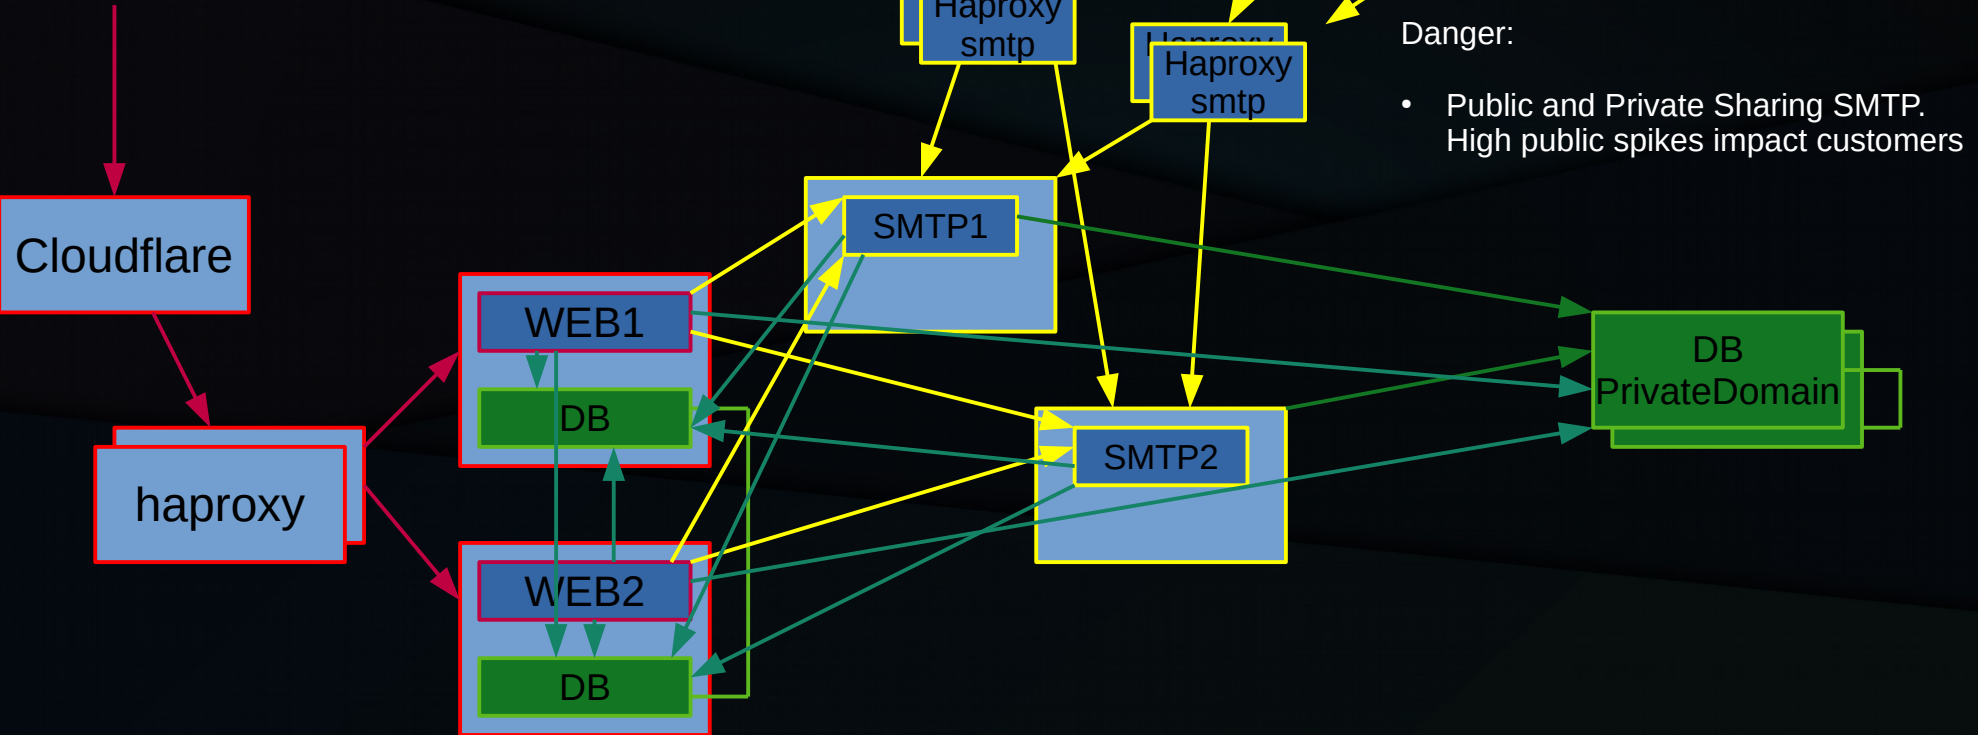
# Private Domains

Public Email  
MX1:25

Public Email  
MX2:25

mail.privatedomain.com  
MX:25

WEB  
HTTPS:443



# Private Domains

Public Email  
MX1:25

Public Email  
MX2:25

mail.privatedomain.com  
MX:25

WEB  
HTTPS:443

Cloudflare

haproxy

WEB1

DB

WEB2

DB

SMTP1

SMTP2

SMTP-P

DB

PrivateDomain

So:

- Add SMTP server and MX record just for private domains
- Scale Private System as growth demands

# Everything is Great!

- Subscribers get their own Private Mailinator
- Public users still get their public system
- Danger: Every incoming email (i.e. thousands per second) must check database to see if THIS domain is a private one
  - Solution: Replicate Private Domain DB in memory

# Big Idea: Websockets

- Currently
  - All clients “poll” the server for updates to the inbox (i.e. has new email arrived?)
  - If we have say, 2000 concurrent users who each poll an inbox every 5 seconds
    - 24,000 web hits a minute. 400 per second.
    - That’s a lot of useless web traffic
- Incoming email is a “stream”. Why not “tap into” that stream by specifying an inbox.
- Think of Mailinator as a “Email Twitter” where you “subscribe” to Inboxes
- Arriving email will INSTANTLY arrive in the web UI
  - Should we write our own pub/sub system? (Of course we should – and we do)
  - But eventually decide Redis on localhost has better features



# RealTime Inboxes !

- Each SMTP server gets it's own Redis and publishes all “inbox snippets” of incoming email to Redis
  - To, From, Subject, Date
  - No email-body
  - No data leaves the SMTP server unless it's subscribed to
  - We can abuse Redis as needed
    - Very robust software
    - Running on localhost – no bandwidth!

# (No) Websockets

Public Email  
MX1:25

Public Email  
MX2:25

mail.privatedomain.com  
MX:25

WEB  
HTTPS:443

Cloudflare

haproxy

WEB1  
DB

WEB2  
DB

SMTP1

SMTP2

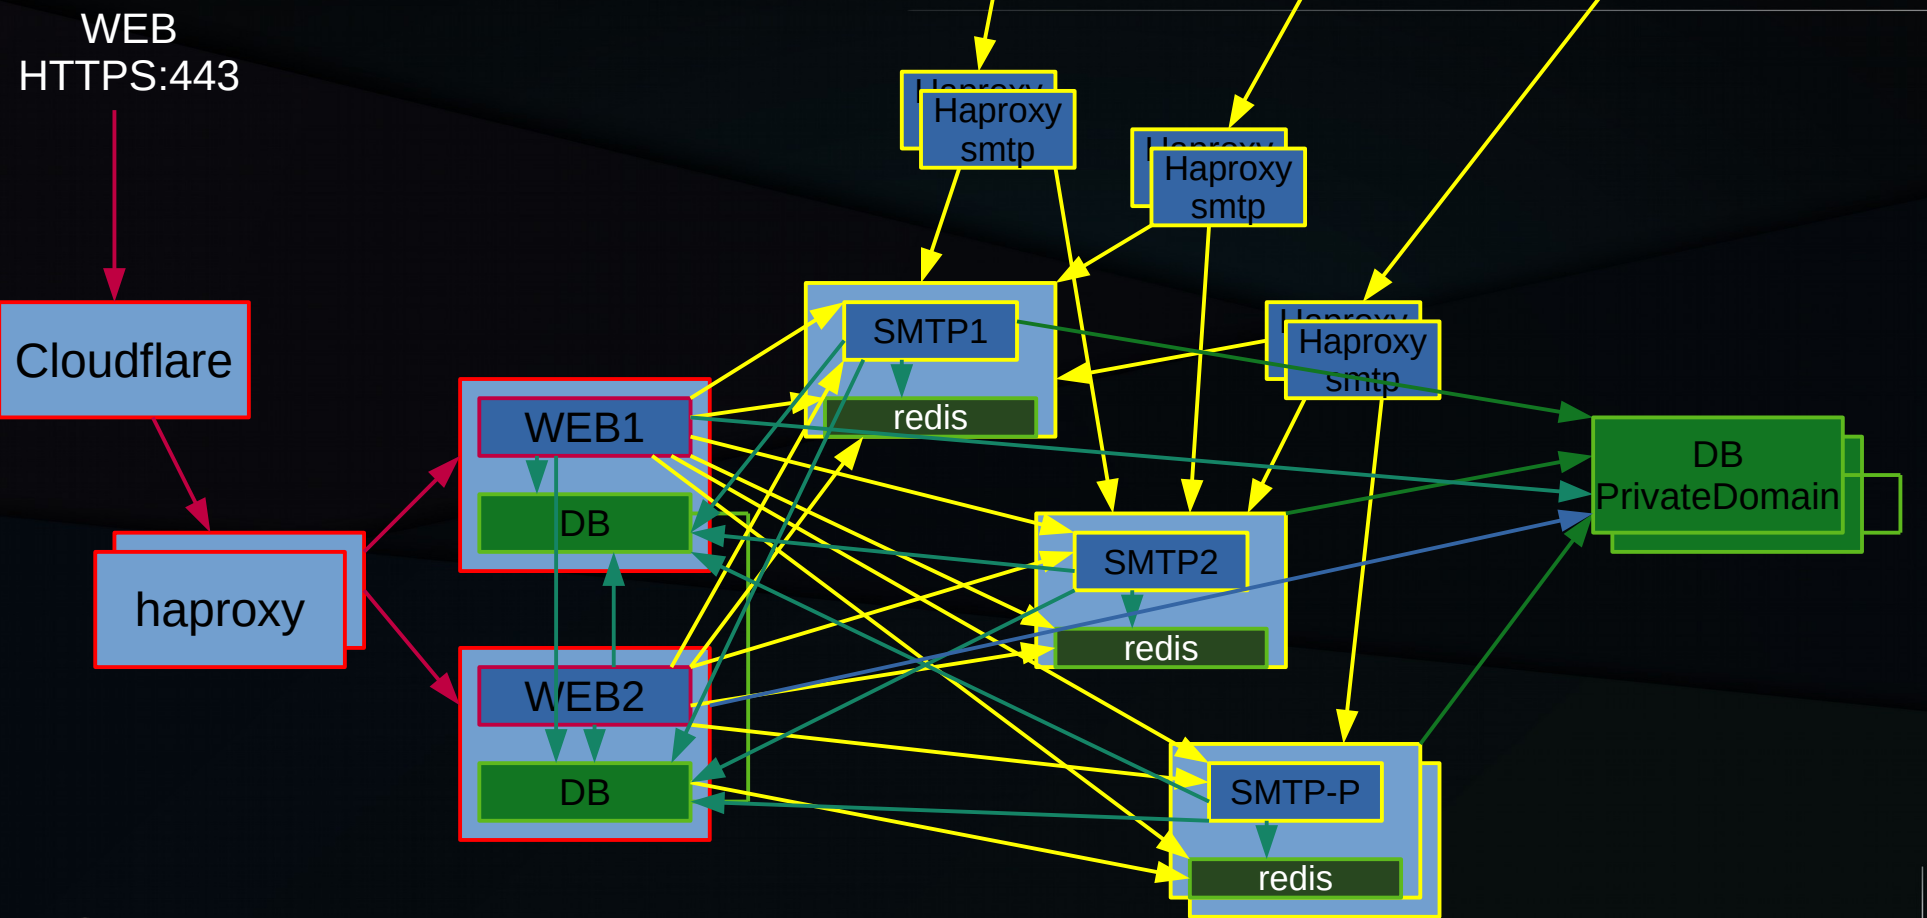
SMTP-P

Haproxy smtp

DB  
PrivateDomain



# Websockets



# Everything is Great!

- I think?
- I mean, it seems like things are running good
- Actually I have no idea what's going on
  - We have zero system monitoring
- Need full system & application monitoring

# ELK

---

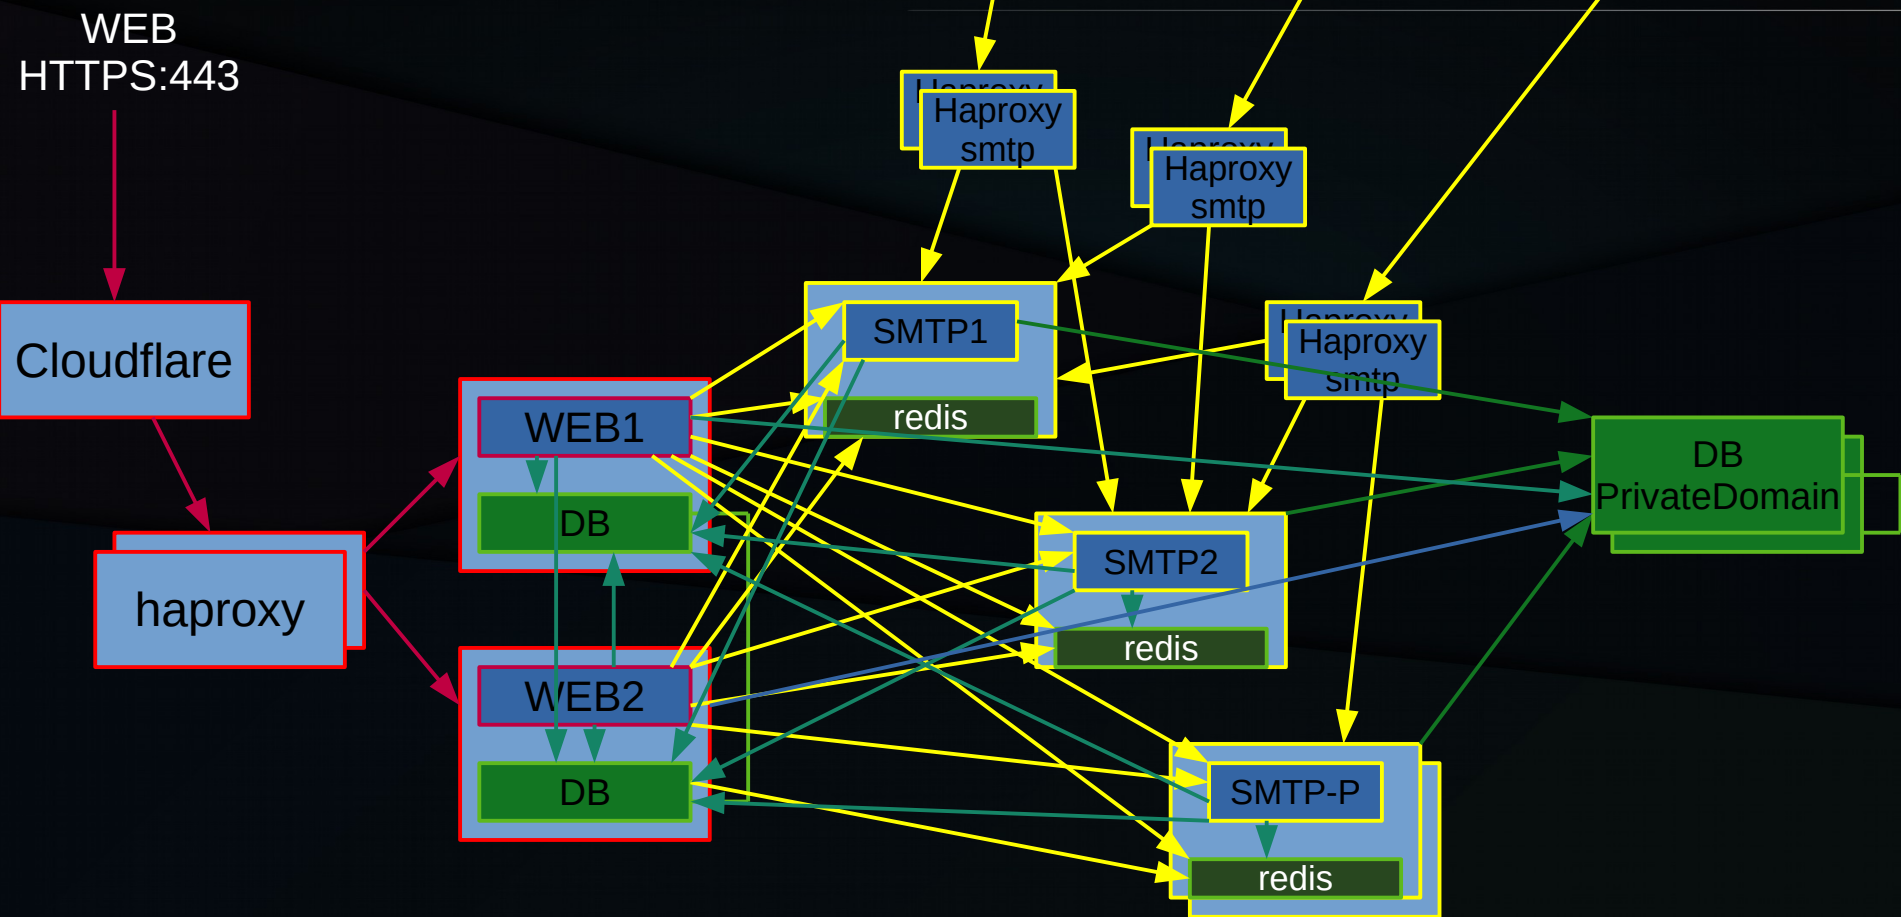
- Elasticsearch
- Logstash
- Kibana (and Grafana)

# ELK-G

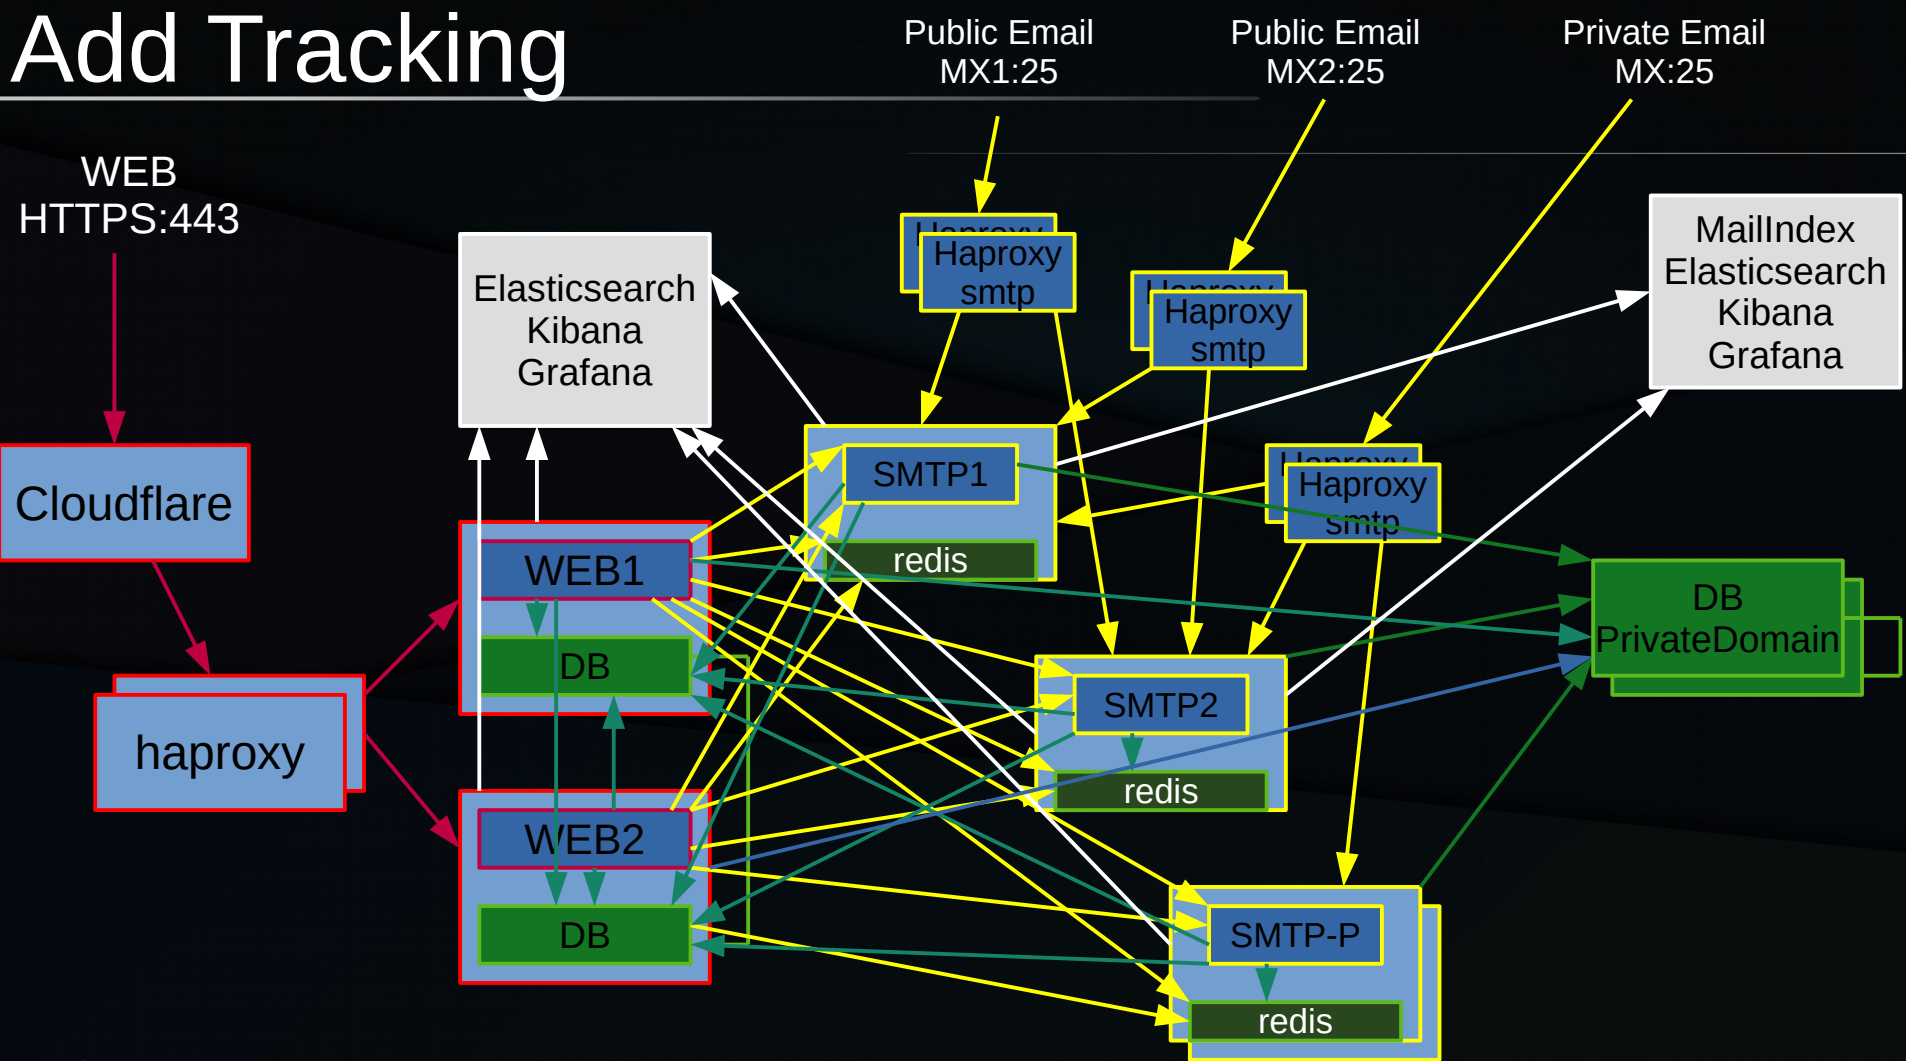
- ElasticSearch
  - A very cool, mostly free search engine for JSON data
- Logstash
  - A system where you take your structured program data, unstructure it into loglines, and logstash will re-structure it for you into JSON
  - That sounds dumb. Let's not use that
- Kibana
  - Search engine front-end. Like Google with cool graphs (Elasticsearch GUI)
- Graphana
  - Like Kibana, but only graphs. Very powerful (another GUI for Elasticsearch)



# (No) Tracking



# Add Tracking



# EKG

---

- Cycle all analytic data every few days
- Let's see the data!

# Mailinator Today

- Team of engineers, sales, marketing, support
- ~3-10MM emails/day
  - This is way down from 20-30MM a few years ago
  - Biggest Public spike seen:
    - ~200MM emails/day
    - 2000-3000 emails/second
      - Very hand wavy – how big of emails? How many per connection? How many SMTP commands?
- ~30,000-60,000 unique users/day
  - >100K when we hit the news



# Mailinator Today

- Thousands of corporate users
- Millions of API hits/day
- All Java – few frameworks. SMTP is raw socket code
- Web: Raw Servlets/JSPs, Angular 1.0 + MailinatorFrontEnd™, Websockets

# Disposable Email Games

(more stuff we learned)



# Banning Mailinator.com emails

- What if they ban you?
  - They did... and how
  - At first – we took it personal and created “alternate domains”
    - Remember? Mailinator accepts email for any domain
    - @notmailinator.com
    - @reallmyemail.com
    - @veryrealemail.com
    - @putthisinyourspamdatabase.com
- “Hi, could I get a list of all your alternate domains?”
  - No.
  - We never showed the entire list
- Sites would scrape us every day to get a list of the alternate domains

# Banning Mailinator.com emails

- Soon however, we realized this wasn't a "war" we wanted to win
- It's your site. If you want to block domains, that's your business
- But note, It's probably not completely possible
  - There's dozens/hundreds of disposable email sites now
  - Here's a list of ~23,000 domains to block to get you going

<https://github.com/ivolo/disposable-email-domains/blob/master/index.json>

# A lesson in defensive programming

- Your boss tells you to ban all @mailinator.com signups, so you write something like:

```
if (email.getDomain().equals("mailinator.com")) {  
    // reject the signup attempt  
}
```

# A lesson in defensive programming

- Remember we said the Mailinator SMTP server accepts email ANY domain?
  - That includes subdomains of mailinator.com
  - So Public Mailinator not only accepts email for
    - <anything>@mailinator.com
  - It also accepts email for:
    - <anything>@<anything>.mailinator.com

# A lesson in defensive programming

```
if (email.getDomain().endsWith("mailinator.com")) {  
    // reject the signup attempt  
}
```

# Healthcare.gov

- Many companies that specialize in signing people up for healthcare.gov use Mailinator addresses
  - i.e. non-technical people may not even have an email address nor know how to get one
  - Spike in November each year
  - Some subscribe for the month for privacy and to keep ownership of domain



# hhhhhhhhhhhhhhhhhhhh@mailinator.com

- May 3, 2017 – someone creates a Google App store application called “Google Docs”
  - Hint: it wasn’t google. It was just called that
- Then emails hundreds of people this:



# hhhhhhhhhhhhhhhhhhhh@mailinator.com

- They then get a message, “Google Docs wants to view your Contacts. Ok?”
- When they click OK, it emails all their contacts the same thing
- It spread across the Internet like wildfire

# hhhhhhhhhhhhhhhhhh@mailinator.com

- Note that when it emailed their contacts, it put their entire contacts list in the BCC field
- But the program needed something in the “to” field.
  - They chose [hhhhhhhhhhhhhhhhhh@mailinator.com](mailto:hhhhhhhhhhhhhhhhhh@mailinator.com)
- That inbox got an email every time the virus spread
  - i.e. It received hundreds of thousands of emails that day

# FIN

